

CYO Project - Hospital death

Ronaldo Benjamin Montes Kohler

25/2/2022

INTRODUCTION :

Machine learning is a big challenge it requires a lot mathematics, skills, time and patience, in this project we are going to analyze and try to predict a hospital death based on different predictors and different algorithms (glm , knn , random forest , rpart), the principal objective is to reduce the Rmse (an statistic formula that measures the amount of error between two sets of data), the data comes from a page called kaggle and the creator and owner of the information you can find it in the following link .

<https://www.kaggle.com/mitishaagarwal/patient>.

Loading and downloading all the packages from R that we are going to use in this project.

```
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")
if(!require(randomForest)) install.packages("randomForest", repos = "http://cran.us.r-project.org")
if(!require(rpart)) install.packages("rpart", repos = "http://cran.us.r-project.org")
if(!require(rpart.plot)) install.packages("rpart.plot", repos = "http://cran.us.r-project.org")
if(!require(ranger)) install.packages("ranger", repos = "http://cran.us.r-project.org")
if(!require(class)) install.packages("class", repos = "http://cran.us.r-project.org")
if(!require(gam)) install.packages("gam", repos = "http://cran.us.r-project.org")
if(!require(splines)) install.packages("splines", repos = "http://cran.us.r-project.org")
if(!require(foreach)) install.packages("foreach", repos = "http://cran.us.r-project.org")
if(!require(MASS)) install.packages("MASS", repos = "http://cran.us.r-project.org")
if(!require(C50)) install.packages("C50", repos = "http://cran.us.r-project.org")

library(tidyverse)
library(caret)
library(data.table)
```

```
library(lubridate)
library(randomForest)
library(dslabs)
library(rpart)
library(rpart.plot)
library(ranger)
library(class)
library(MASS)
library(gam)
library(splines)
library(foreach)
library(C50)
```

Methods/analysis :

Downloading the csv data and reading into a data frame.

```
Survival_Prediction <-read.csv('https://www.dropbox.com/s/angh53kkzc4awo6/dataset.csv?dl=1')
```

Now first we look at the data from the package with the function `head()`, it will illustrate what it contains the columns names and the type of data inside.

```
head(Survival_Prediction)
```

```
##   encounter_id patient_id hospital_id age   bmi elective_surgery ethnicity
## 1      66154      25312         118  68 22.73                0 Caucasian
## 2     114252     59342          81  77 27.42                0 Caucasian
## 3     119783     50777         118  25 31.95                0 Caucasian
## 4      79267     46918         118  81 22.64                1 Caucasian
## 5      92056     34377          33  19  NA                0 Caucasian
## 6      33181     74489          83  67 27.56                0 Caucasian
##   gender height          icu_admit_source icu_id icu_stay_type    icu_type
## 1     M  180.3                Floor      92      admit      CTICU
## 2     F  160.0                Floor      90      admit Med-Surg ICU
## 3     F  172.7 Accident & Emergency      93      admit Med-Surg ICU
## 4     F  165.1 Operating Room / Recovery      92      admit      CTICU
## 5     M  188.0 Accident & Emergency      91      admit Med-Surg ICU
## 6     M  190.5 Accident & Emergency      95      admit Med-Surg ICU
##   pre_icu_los_days weight apache_2_diagnosis apache_3j_diagnosis
## 1    0.541666667    73.9             113          502.01
## 2    0.927777778    70.2             108          203.01
## 3    0.000694444    95.3             122          703.03
## 4    0.000694444    61.7             203         1206.03
## 5    0.073611111     NA             119          601.01
```

```

## 6      0.000694444 100.0      301      403.01
##  apache_post_operative arf_apache gcs_eyes_apache gcs_motor_apache
## 1           0           0           3           6
## 2           0           0           1           3
## 3           0           0           3           6
## 4           1           0           4           6
## 5           0           0           NA          NA
## 6           0           0           4           6
##  gcs_unable_apache gcs_verbal_apache heart_rate_apache intubated_apache
## 1           0           4          118           0
## 2           0           1          120           0
## 3           0           5          102           0
## 4           0           5          114           1
## 5          NA          NA           60           0
## 6           0           5          113           0
##  map_apache resprate_apache temp_apache ventilated_apache d1_diasbp_max
## 1          40          36          39.3           0          68
## 2          46          33          35.1           1          95
## 3          68          37          36.7           0          88
## 4          60           4          34.8           1          48
## 5         103          16          36.7           0          99
## 6         130          35          36.6           0         100
##  d1_diasbp_min d1_diasbp_noninvasive_max d1_diasbp_noninvasive_min
## 1           37           68           37
## 2           31           95           31
## 3           48           88           48
## 4           42           48           42
## 5           57           99           57
## 6           61          100           61
##  d1_heartrate_max d1_heartrate_min d1_mbp_max d1_mbp_min
## 1          119           72           89           46
## 2          118           72          120           38
## 3           96           68          102           68
## 4          116           92           84           84
## 5           89           60          104           90
## 6          113           83          127           80
##  d1_mbp_noninvasive_max d1_mbp_noninvasive_min d1_resprate_max d1_resprate_min
## 1           89           46           34           10
## 2          120           38           32           12
## 3          102           68           21            8
## 4           84           84           23            7
## 5          104           90           18           16
## 6          127           80           32           10
##  d1_spo2_max d1_spo2_min d1_sysbp_max d1_sysbp_min d1_sysbp_noninvasive_max
## 1          100           74          131           73          131
## 2          100           70          159           67          159
## 3           98           91          148          105          148
## 4          100           95          158           84          158
## 5          100           96          147          120          147
## 6           97           91          173          107          173
##  d1_sysbp_noninvasive_min d1_temp_max d1_temp_min h1_diasbp_max h1_diasbp_min
## 1           73          39.9          37.2           68           63
## 2           67          36.3          35.1           61           48
## 3          105          37.0          36.7           88           58

```

## 4	84	38.0	34.8	62	44
## 5	120	37.2	36.7	99	68
## 6	107	36.8	36.6	89	89
##	h1_diasbp_noninvasive_max	h1_diasbp_noninvasive_min	h1_hearttrate_max		
## 1	68		63	119	
## 2	61		48	114	
## 3	88		58	96	
## 4	NA		NA	100	
## 5	99		68	89	
## 6	89		89	83	
##	h1_hearttrate_min	h1_mbp_max	h1_mbp_min	h1_mbp_noninvasive_max	
## 1	108	86	85	86	
## 2	100	85	57	85	
## 3	78	91	83	91	
## 4	96	92	71	NA	
## 5	76	104	92	104	
## 6	83	111	111	111	
##	h1_mbp_noninvasive_min	h1_resprate_max	h1_resprate_min	h1_spo2_max	
## 1	85	26	18	100	
## 2	57	31	28	95	
## 3	83	20	16	98	
## 4	NA	12	11	100	
## 5	92	NA	NA	100	
## 6	111	12	12	97	
##	h1_spo2_min	h1_sysbp_max	h1_sysbp_min	h1_sysbp_noninvasive_max	
## 1	74	131	115	131	
## 2	70	95	71	95	
## 3	91	148	124	148	
## 4	99	136	106	NA	
## 5	100	130	120	130	
## 6	97	143	143	143	
##	h1_sysbp_noninvasive_min	d1_glucose_max	d1_glucose_min	d1_potassium_max	
## 1	115	168	109	4.0	
## 2	71	145	128	4.2	
## 3	124	NA	NA	NA	
## 4	NA	185	88	5.0	
## 5	120	NA	NA	NA	
## 6	143	156	125	3.9	
##	d1_potassium_min	apache_4a_hospital_death_prob	apache_4a_icu_death_prob	aids	
## 1	3.4		0.10	0.05	0
## 2	3.8		0.47	0.29	0
## 3	NA		0.00	0.00	0
## 4	3.5		0.04	0.03	0
## 5	NA		NA	NA	0
## 6	3.7		0.05	0.02	0
##	cirrhosis	diabetes_mellitus	hepatic_failure	immunosuppression	leukemia
## 1	0	1	0	0	0
## 2	0	1	0	0	0
## 3	0	0	0	0	0
## 4	0	0	0	0	0
## 5	0	0	0	0	0
## 6	0	1	0	0	0
##	lymphoma	solid_tumor_with_metastasis	apache_3j_bodysystem	apache_2_bodysystem	
## 1	0	0	Sepsis	Cardiovascular	

```
## 2      0      0      Respiratory      Respiratory
## 3      0      0      Metabolic        Metabolic
## 4      0      0      Cardiovascular    Cardiovascular
## 5      0      0      Trauma            Trauma
## 6      0      0      Neurological      Neurologic
##      X hospital_death
## 1 NA      0
## 2 NA      0
## 3 NA      0
## 4 NA      0
## 5 NA      0
## 6 NA      0
```

We notice that there is a column called X in the data frame which most of the data is NA, we clean the data from any NA value.

```
Survival_Prediction <- Survival_Prediction[,-84]

Survival_Prediction <- Survival_Prediction %>%
  drop_na()

sum(is.na(Survival_Prediction$age))
```

```
## [1] 0
```

We also notice that there is a lot of specificity in the height , weight and the temperature, It will not be very helpful to have the data like this since when grouping them there would be many groups of a single value and the analysis would be much more difficult, we round the values to the nearest integer.

```
Survival_Prediction$height <- round(Survival_Prediction$height)
Survival_Prediction$weight <- round(Survival_Prediction$weight)
Survival_Prediction$temp_apache <- round(Survival_Prediction$temp_apache)
Survival_Prediction$d1_temp_min <- round(Survival_Prediction$d1_temp_min )
Survival_Prediction$d1_temp_max <- round(Survival_Prediction$d1_temp_max )
```

Due to data visualization problems we have decided to create a special variable.

```
Survival_Prediction_vi <- Survival_Prediction
```

Some columns has binary data or specific groups, so we decide to convert it to a factor variable.

```
Survival_Prediction$hospital_death <- as.factor(Survival_Prediction$hospital_death)
Survival_Prediction$gcs_eyes_apache <- as.factor(Survival_Prediction$gcs_eyes_apache)
Survival_Prediction$gcs_motor_apache <- as.factor(Survival_Prediction$gcs_motor_apache)
Survival_Prediction$gcs_verbal_apache <- as.factor(Survival_Prediction$gcs_verbal_apache)
Survival_Prediction$intubated_apache <- as.factor(Survival_Prediction$intubated_apache)
Survival_Prediction$aids <- as.factor(Survival_Prediction$aids)
Survival_Prediction$cirrhosis <- as.factor(Survival_Prediction$cirrhosis)
Survival_Prediction$solid_tumor_with_metastasis <- as.factor(Survival_Prediction$solid_tumor_with_metastasis)
```

We are going to split the data in two parts one for training the algorithm (90%), and another to validate(10%) the results.

```
test_index <- createDataPartition(y = Survival_Prediction$hospital_death, times = 1, p = 0.1, list = FALSE)
edx <- Survival_Prediction[-test_index,]
temp <- Survival_Prediction[test_index,]

validation <- temp %>%
  semi_join(edx, by = "hospital_death")

removed <- anti_join(temp, validation)
```

```
## Joining, by = c("encounter_id", "patient_id", "hospital_id", "age", "bmi", "elective_surgery", "ethnicity")
```

```
edx <- rbind(edx, removed)
```

We are only going to use the EDX set to realize all the machine learning work, so the EDX data set is going to be splitted on a train set(90%) and a test set(10%).

```
test_index <- createDataPartition(y = edx$hospital_death, times = 1, p = 0.1, list = FALSE)
edx_train <- edx[-test_index,]
temp <- edx[test_index,]
```

```

edx_test <- temp %>%
  semi_join(edx, by = "hospital_death")

removed <- anti_join(temp, edx_test)

## Joining, by = c("encounter_id", "patient_id", "hospital_id", "age", "bmi", "elective_surgery", "ethn

edx_train <- rbind(edx_train, removed)

```

The data has 84 columns, trying to explore all of these columns might take a lot of time, so we are only going to use a few columns more flashly.

```
ncol(edx_train)
```

```
## [1] 84
```

Some columns have the label Apache that is an acronym for Acute Physiology and Chronic Health Evaluation, It is applied within 24 hours of admission of a patient to an intensive care unit (ICU).

```
colnames(edx_train)
```

```

## [1] "encounter_id"      "patient_id"
## [3] "hospital_id"       "age"
## [5] "bmi"               "elective_surgery"
## [7] "ethnicity"         "gender"
## [9] "height"            "icu_admit_source"
## [11] "icu_id"            "icu_stay_type"
## [13] "icu_type"          "pre_icu_los_days"
## [15] "weight"            "apache_2_diagnosis"
## [17] "apache_3j_diagnosis" "apache_post_operative"
## [19] "arf_apache"        "gcs_eyes_apache"
## [21] "gcs_motor_apache"  "gcs_unable_apache"
## [23] "gcs_verbal_apache" "heart_rate_apache"
## [25] "intubated_apache"  "map_apache"
## [27] "resprate_apache"   "temp_apache"
## [29] "ventilated_apache" "d1_diasbp_max"
## [31] "d1_diasbp_min"     "d1_diasbp_noninvasive_max"
## [33] "d1_diasbp_noninvasive_min" "d1_heartrate_max"
## [35] "d1_heartrate_min"  "d1_mbp_max"
## [37] "d1_mbp_min"        "d1_mbp_noninvasive_max"
## [39] "d1_mbp_noninvasive_min" "d1_resprate_max"
## [41] "d1_resprate_min"   "d1_spo2_max"

```

## [43] "d1_spo2_min"	"d1_sysbp_max"
## [45] "d1_sysbp_min"	"d1_sysbp_noninvasive_max"
## [47] "d1_sysbp_noninvasive_min"	"d1_temp_max"
## [49] "d1_temp_min"	"h1_diasbp_max"
## [51] "h1_diasbp_min"	"h1_diasbp_noninvasive_max"
## [53] "h1_diasbp_noninvasive_min"	"h1_heartrate_max"
## [55] "h1_heartrate_min"	"h1_mbp_max"
## [57] "h1_mbp_min"	"h1_mbp_noninvasive_max"
## [59] "h1_mbp_noninvasive_min"	"h1_resprate_max"
## [61] "h1_resprate_min"	"h1_spo2_max"
## [63] "h1_spo2_min"	"h1_sysbp_max"
## [65] "h1_sysbp_min"	"h1_sysbp_noninvasive_max"
## [67] "h1_sysbp_noninvasive_min"	"d1_glucose_max"
## [69] "d1_glucose_min"	"d1_potassium_max"
## [71] "d1_potassium_min"	"apache_4a_hospital_death_prob"
## [73] "apache_4a_icu_death_prob"	"aids"
## [75] "cirrhosis"	"diabetes_mellitus"
## [77] "hepatic_failure"	"immunosuppression"
## [79] "leukemia"	"lymphoma"
## [81] "solid_tumor_with_metastasis"	"apache_3j_bodysystem"
## [83] "apache_2_bodysystem"	"hospital_death"

In this project we are going to use some machine learning algorithms, to analyze and use the algorithms we need to understand how they work.

Random forest is a supervised machine learning algorithm that uses ensemble learning method for regression, ‘the idea of random forests is to generate many predictors, each using regression or classification trees, and then forming a final prediction based on the average prediction of all these trees’.

C5.0 model works by splitting the sample based on the field that provides the maximum information gain. Each sub-sample defined by the first split is then split again, usually based on a different field, and the process repeats until the sub samples cannot be split any further. Finally, the lowest-level splits are reexamined.

Generalized Additive Model using LOESS

The additive model generalizes the linear model by modeling the expected value of Y as

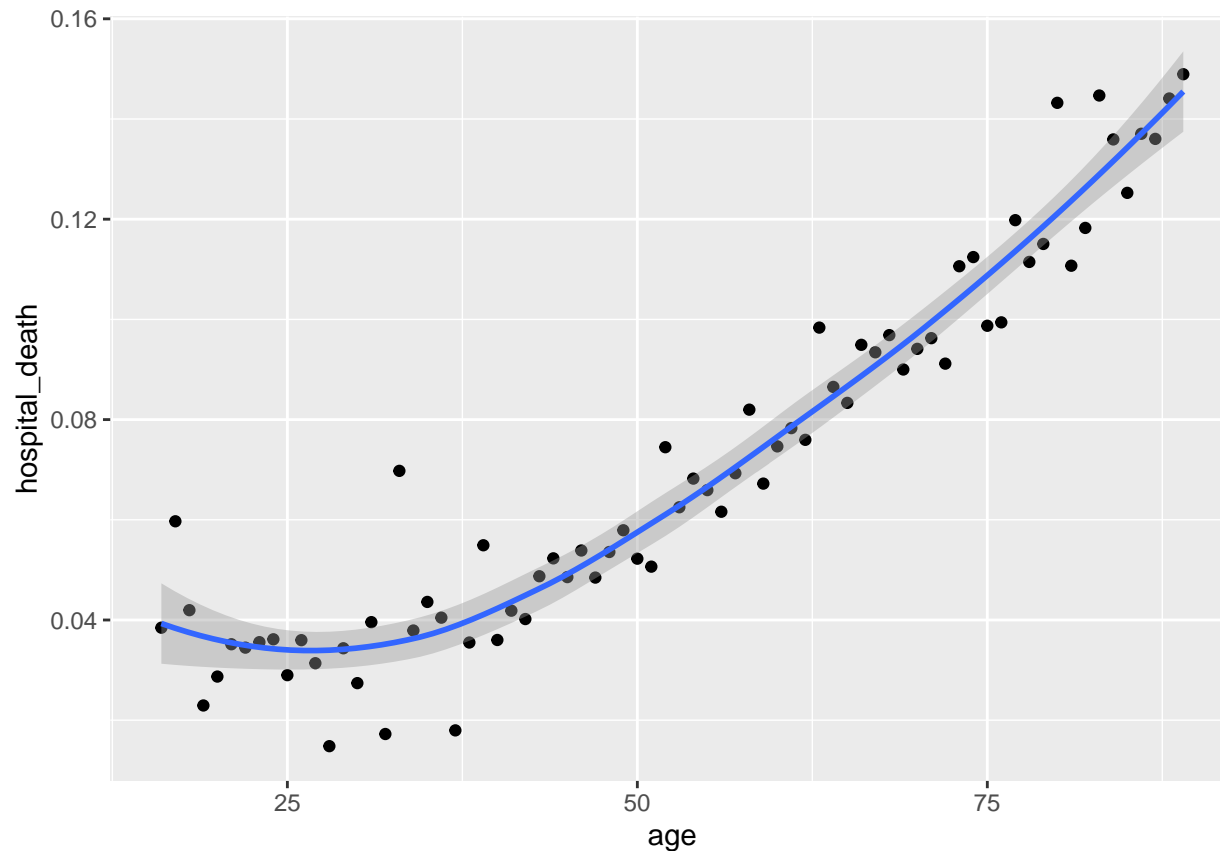
$$E(Y) = f(X_1 \dots X_p) = s_0 + s_1(X_1) + \dots + s_p(X_p)$$

Where $s_i(X)$, $i = 1 \dots p$, are smooth functions. These functions are estimated in a nonparametric fashion.

Exploring the data the first intuition that we had was that the probability of death of a older person was greater than a young, so we plot the mean hospital death grouped by the age.

```
Survival_Prediction_vi %>%
  group_by(age) %>%
  summarize(hospital_death = mean(hospital_death), age = mean(age)) %>%
  ggplot(aes(age, hospital_death)) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



There's a strong evidence that an older person has a higher of die than a younger one.

RandomForest() function allows us to run Random Forest Regression.

Testing the age on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age, data = edx_train)
y_hat_rf <- predict(fit_rf, edx_test)
mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.914159
```

```
sum(y_hat_rf == 1)
```

```
## [1] 0
```

`Train()` function with `method = 'C5.0Tree'` parameter allows us to run C5.0Tree algorithm.

Testing the first model C5.0Tree.

```
fit_ct <- train(hospital_death ~ age, data = edx_train, method = 'C5.0Tree' )
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.914159
```

```
sum(y_hat_ct == 1)
```

```
## [1] 0
```

`Train()` function with `method = 'gamLoess'` parameter allows us to run gamLoess algorithm.

Testing the first model gamLoess.

```
fit_gam <- train(hospital_death ~ age, data = edx_train, method = 'gamLoess' )
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

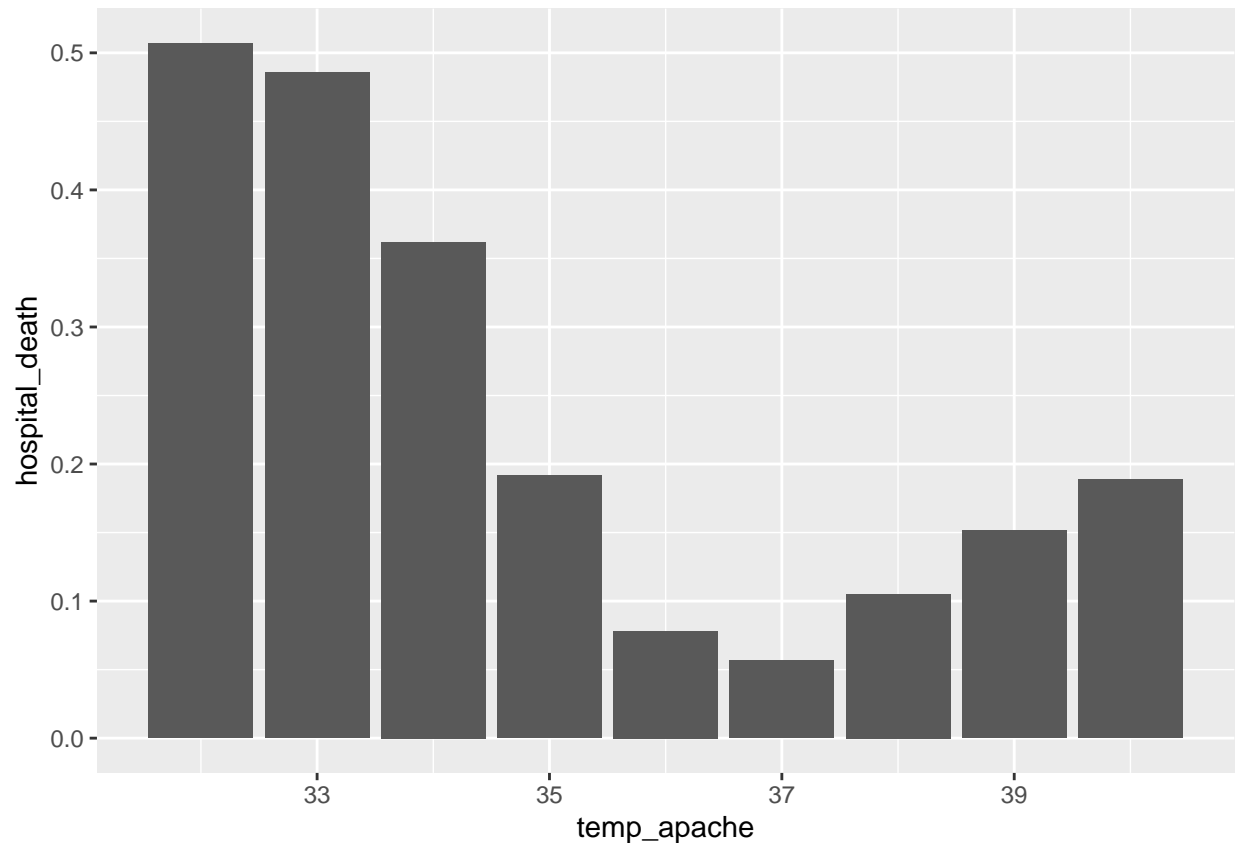
```
## [1] 0.914159
```

```
sum(y_hat_gam == 1)
```

```
## [1] 0
```

The person's temperature can be a great indication of risk of death especially when entering intensive care.

```
Survival_Prediction_vi %>%
  group_by(temp_apache ) %>%
  summarize(temp_apache = mean(temp_apache , na.rm = TRUE), hospital_death = mean(hospital_death , na.rm = TRUE))
ggplot(aes(temp_apache , hospital_death )) +
  geom_col()
```



Testing the temp_apache on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + temp_apache , data = edx_train)
```

```
y_hat_rf <- predict(fit_rf, edx_test)
```

```
mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.913966
```

```
sum(y_hat_rf == 1)
```

```
## [1] 31
```

Testing the temp_apache on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache , data = edx_train , method = 'C5.0Tree')
```

```
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.914159
```

```
sum(y_hat_ct == 1)
```

```
## [1] 0
```

Testing the temp_apache on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache , data = edx_train, method = 'gamLoess' )
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

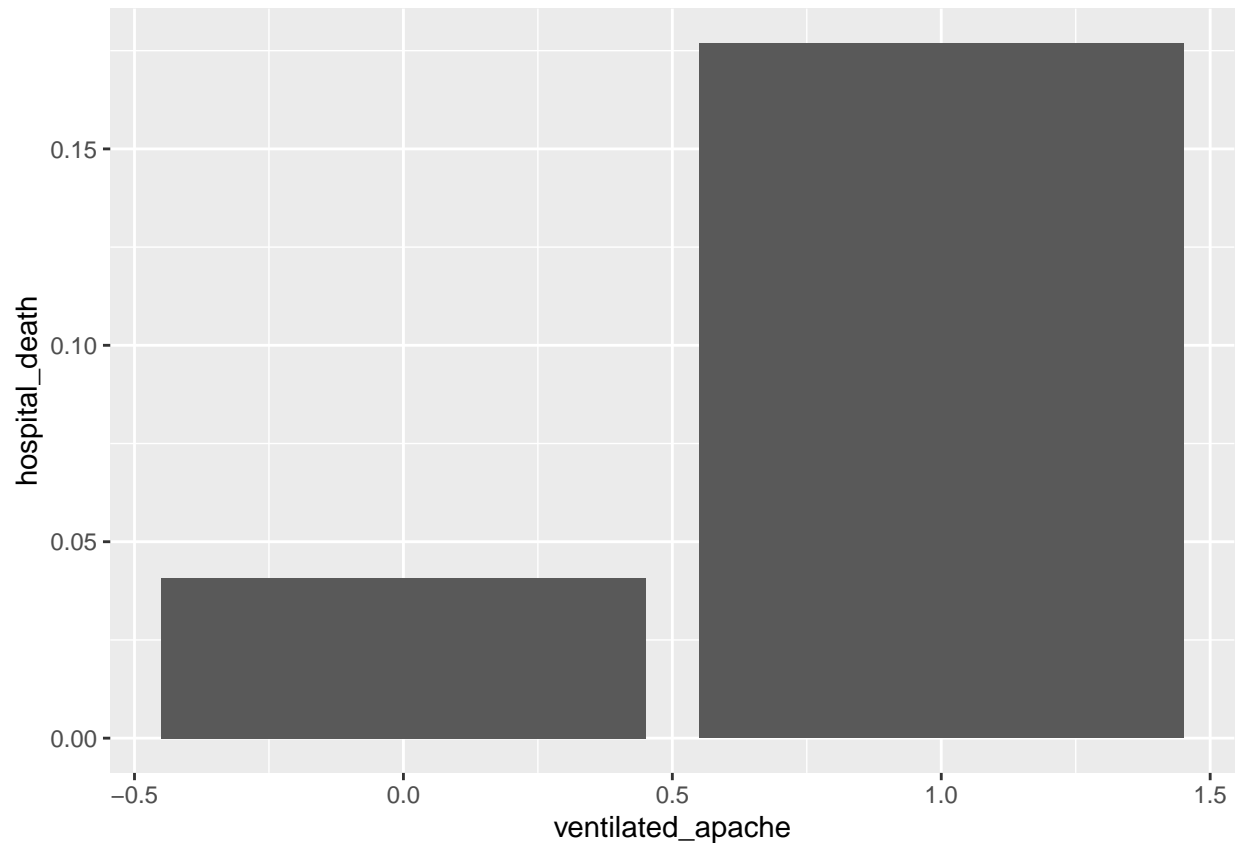
```
## [1] 0.9151235
```

```
sum(y_hat_gam == 1)
```

```
## [1] 7
```

Our intuition says that people who enter the mechanical ventilation area are more likely to die.

```
Survival_Prediction_vi %>%
  group_by(ventilated_apache ) %>%
  summarize(ventilated_apache = mean(ventilated_apache , na.rm = TRUE),hospital_death = mean(hospital_
  ggplot(aes(ventilated_apache ,hospital_death )) +
  geom_col()
```



The data is binary, where 1 one is this with mechanical ventilation and 0 is the opposite Analyzing the graph, we realize that people who actually enter mechanical ventilation have a greater chance of dying.

Testing the ventilated__apache on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache , data = edx_train)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9149306
```

```
sum(y_hat_rf == 1)
```

```
## [1] 30
```

Testing the ventilated__apache on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + ventilated_apache + temp_apache , data = edx_train, method = 'C')
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.914159
```

```
sum(y_hat_ct == 1)
```

```
## [1] 0
```

Testing the ventilated_apache on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache , data = edx_train, method = 'gamLoess')
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

```
## [1] 0.9153164
```

```
sum(y_hat_gam == 1)
```

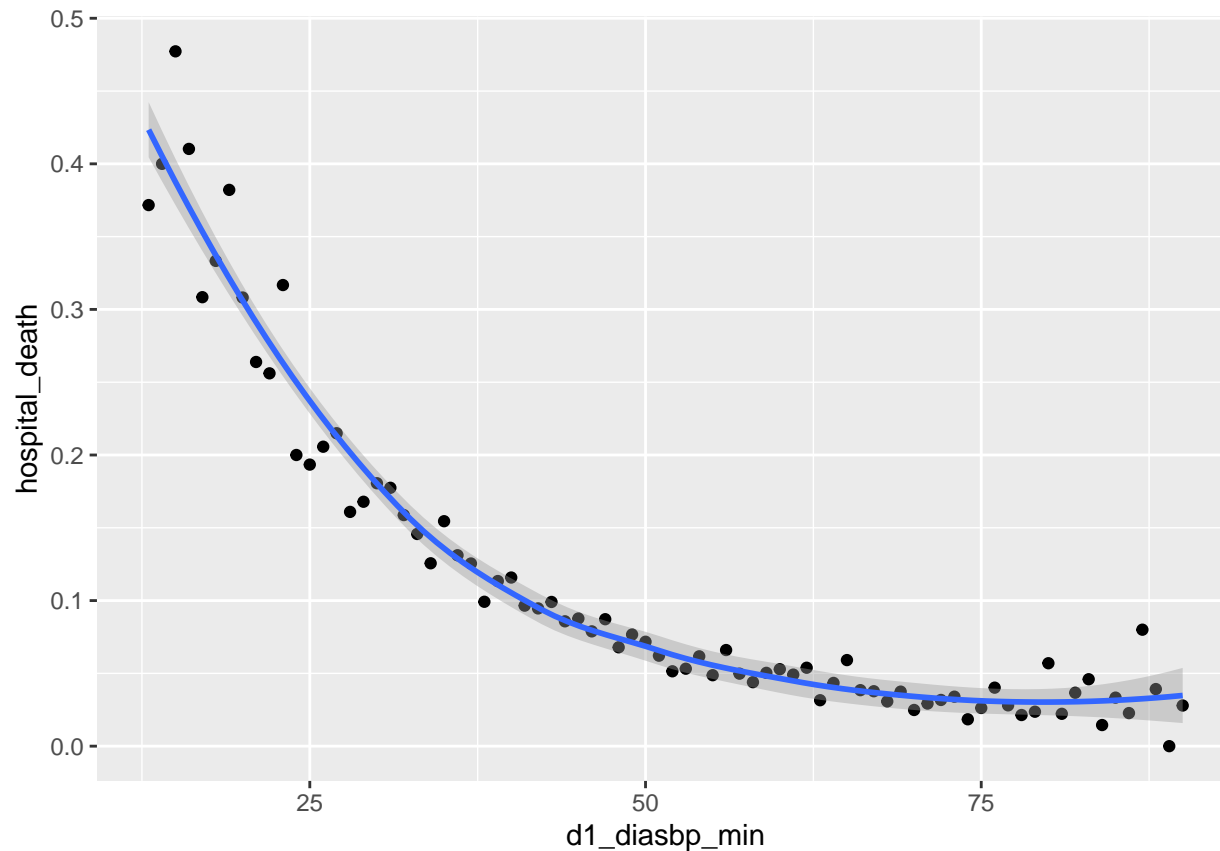
```
## [1] 12
```

Diasbp is an acronym of Diastolic blood pressure.

Can the diastolic blood pressure have an effect on the death of a person?

```
Survival_Prediction_vi %>%
  group_by(d1_diasbp_min) %>%
  summarize(d1_diasbp_min = mean(d1_diasbp_min , na.rm = TRUE), hospital_death = mean(hospital_death))
ggplot(aes(d1_diasbp_min , hospital_death )) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



There is a very clear relationship where having a lower pressure means more chances of death.

Testing the d1_diasbp_min on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min, data = edx_test)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)

## [1] 0.9143519

sum(y_hat_rf == 1)

## [1] 75
```

Testing the d1_diasbp_min on the C5.0Tree algorithm.


```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min , data = edx)
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9151235
```

```
sum(y_hat_ct == 1)
```

```
## [1] 27
```

Testing the d1__diasbp__min on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min , data = edx)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

```
## [1] 0.9157022
```

```
sum(y_hat_gam == 1)
```

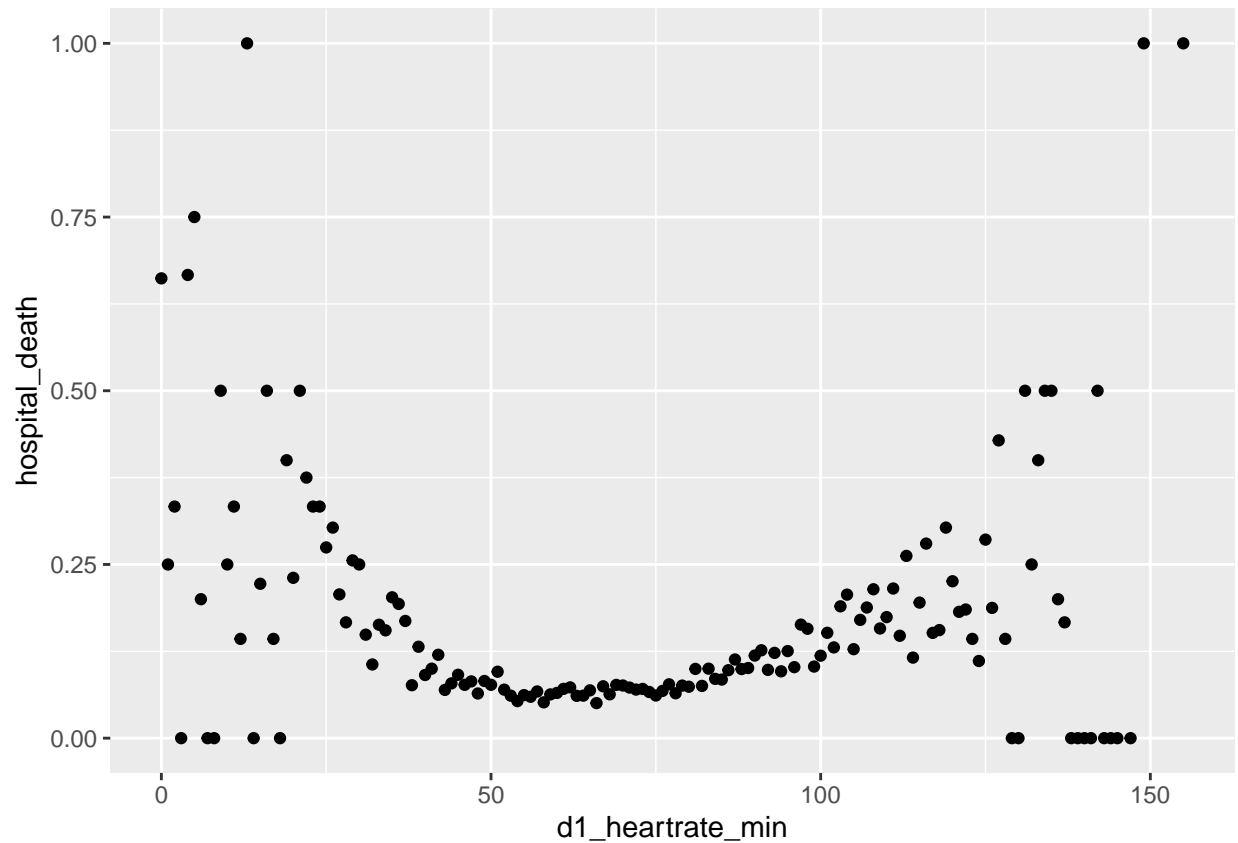
```
## [1] 42
```

Testing the min heart rate on the rpart algorithm.

The heart rate is always a determining factor in the death of a person.

```
Survival_Prediction_vi %>%
  group_by(d1_hearttrate_min ) %>%
  summarize(d1_hearttrate_min = mean(d1_hearttrate_min ), hospital_death = mean(hospital_death , na.rm = TRUE))
ggplot(aes(d1_hearttrate_min , hospital_death , group = d1_hearttrate_min )) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



Testing the `d1_heartrate_min` on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heartrate_min, edx_train)
y_hat_rf <- predict(fit_rf, edx_test)
mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9174383
```

```
sum(y_hat_rf == 1)
```

```
## [1] 101
```

Testing the `d1_heartrate_min` on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_min, edx_train)
y_hat_ct <- predict(fit_ct, edx_test)
```

```
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9178241
```

```
sum(y_hat_ct == 1)
```

```
## [1] 45
```

Testing the `d1_heartrate_min` on the `gamLoess` algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_min,
  data = edx_train)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

```
## [1] 0.9184028
```

```
sum(y_hat_gam == 1)
```

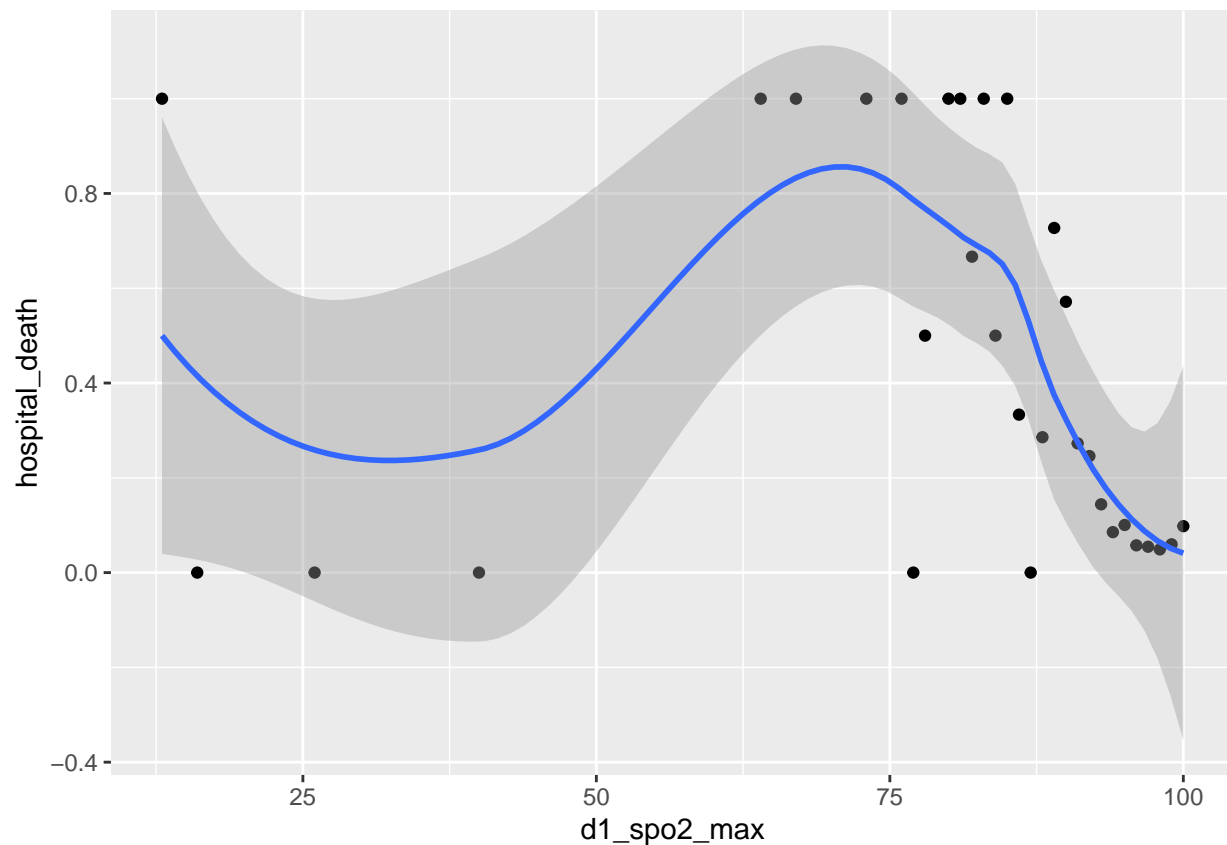
```
## [1] 66
```

SpO2 stands for Saturation of peripheral Oxygen, used to estimate the oxygen saturation of arterial blood .

Can max oxygen saturation of arterial blood have an impact on in-hospital death?

```
Survival_Prediction_vi %>%
  group_by(d1_spo2_max) %>%
  summarize(d1_spo2_max = mean(d1_spo2_max), hospital_death = mean(hospital_death, na.rm = TRUE)) %>%
  ggplot(aes(d1_spo2_max, hospital_death)) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



Testing the d1_spo2_max on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9207176
```

```
sum(y_hat_rf == 1)
```

```
## [1] 94
```

Testing the d1_spo2_max on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heart_rate)
```

```
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9203318
```

```
sum(y_hat_ct == 1)
```

```
## [1] 70
```

Testing the d1_spo2_max on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_max)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

```
## [1] 0.9187886
```

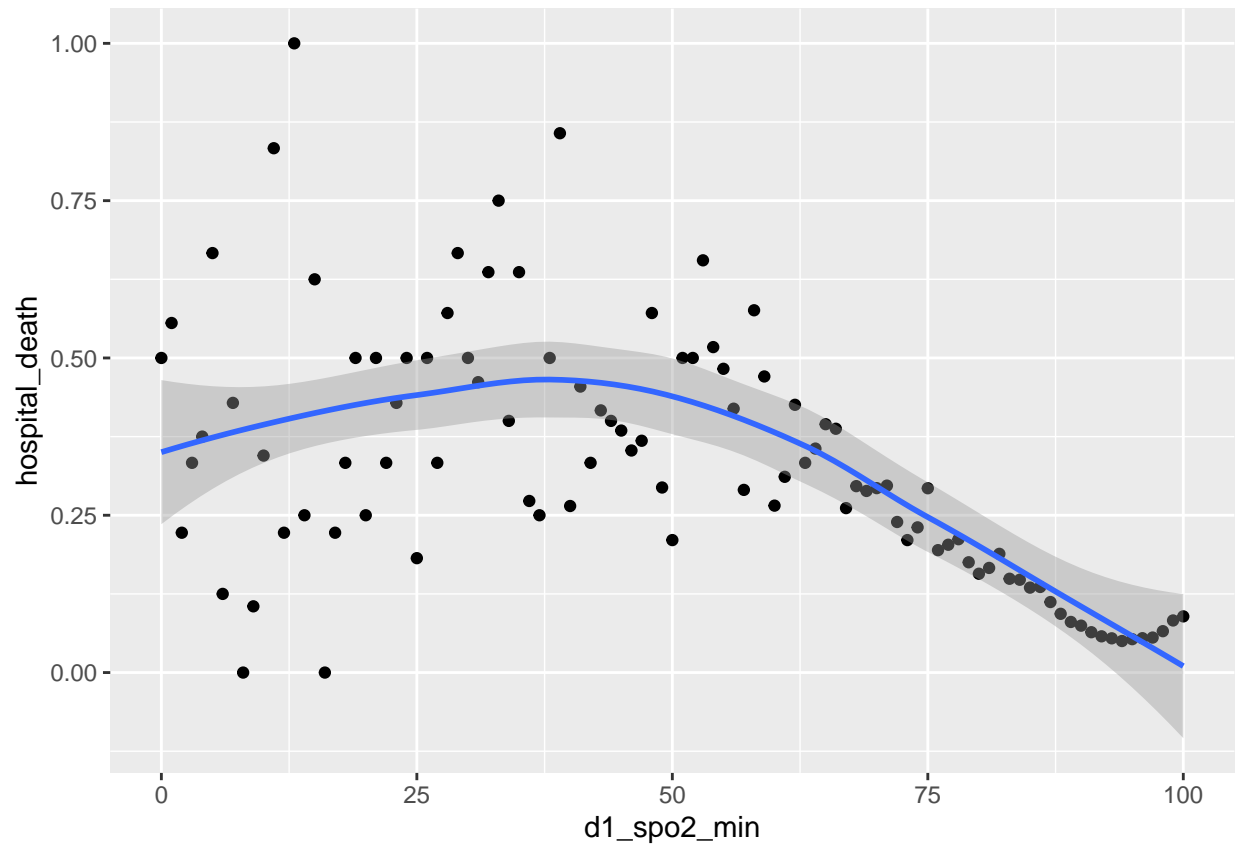
```
sum(y_hat_gam == 1)
```

```
## [1] 68
```

Can min oxygen saturation have an impact on in-hospital death?.

```
Survival_Prediction_vi %>%
  group_by(d1_spo2_min) %>%
  summarize(d1_spo2_min = mean(d1_spo2_min), hospital_death = mean(hospital_death, na.rm = TRUE)) %>%
  ggplot(aes(d1_spo2_min, hospital_death)) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



There is a very strong relationship between the lower the oxygen saturation, the higher the chances of dying

Testing the d1_spo2_min on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate)
y_hat_rf <- predict(fit_rf, edx_test)
mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9207176
```

```
sum(y_hat_rf == 1)
```

```
## [1] 100
```

Testing the d1_spo2_min on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heart_rate)
```

```
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9195602
```

```
sum(y_hat_ct == 1)
```

```
## [1] 76
```

Testing the d1_spo2_min on the gamLoess algorithm.m.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_min)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

```
## [1] 0.9214892
```

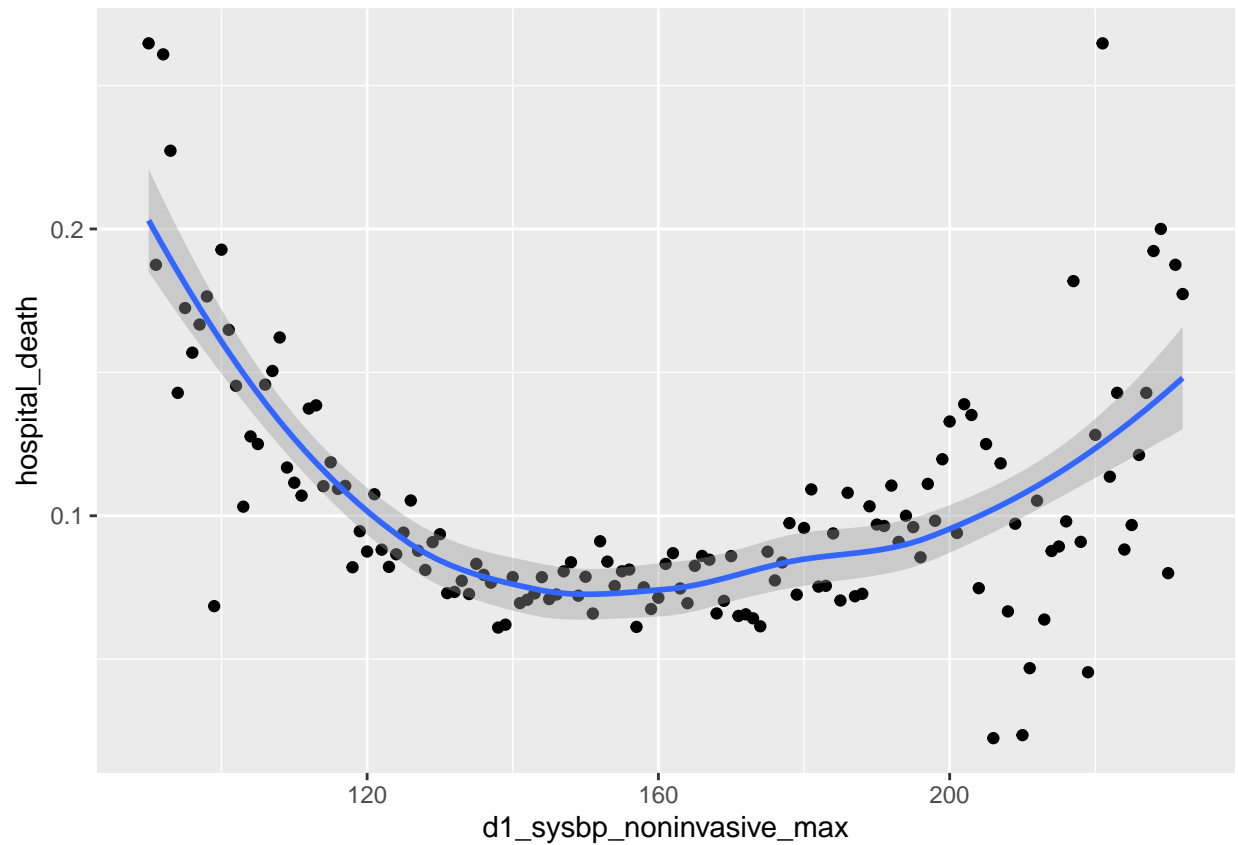
```
sum(y_hat_gam == 1)
```

```
## [1] 88
```

Sysbp or Systolic blood pressures, is the pressure exerted when the heart beats and blood is ejected into the arteries.

```
Survival_Prediction_vi %>%
  group_by(d1_sysbp_noninvasive_max) %>%
  summarize(d1_sysbp_noninvasive_max = mean(d1_sysbp_noninvasive_max), hospital_death = mean(hospital_death)) +
  ggplot(aes(d1_sysbp_noninvasive_max, hospital_death)) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



Testing the `d1_sysbp_noninvasive_max` on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate,
  data = edx_train,
  ntree = 1000,
  importance = TRUE)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9222608
```

```
sum(y_hat_rf == 1)
```

```
## [1] 98
```


Testing the `d1_sysbp_noninvasive_max` on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_max)
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9191744
```

```
sum(y_hat_ct == 1)
```

```
## [1] 80
```

Testing the `d1_sysbp_noninvasive_max` on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_max)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

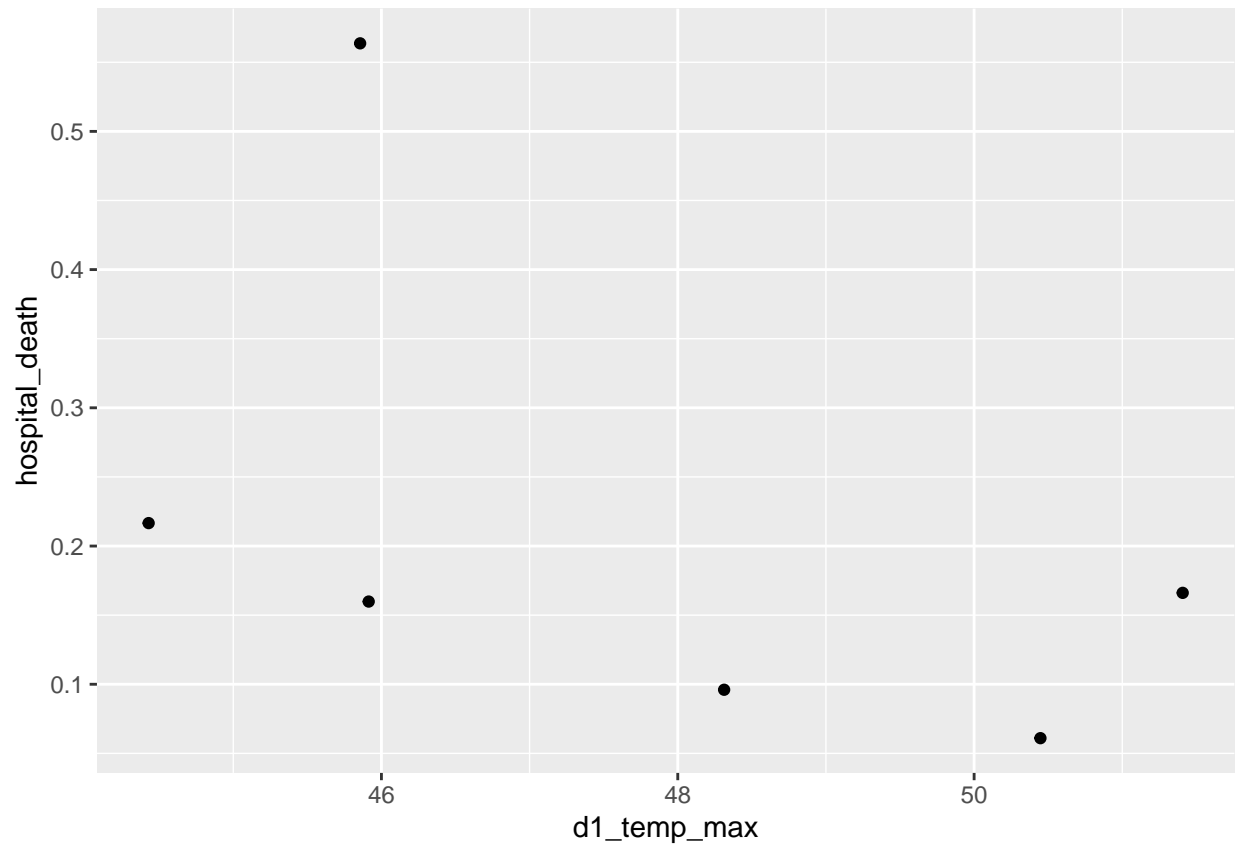
```
## [1] 0.9212963
```

```
sum(y_hat_gam == 1)
```

```
## [1] 93
```

The temperature of people is a great factor to analyze to detect diseases, could this parameter be fatal?

```
Survival_Prediction_vi %>%
  group_by(d1_temp_max) %>%
  summarize(d1_temp_max = mean(d1_diasbp_min), hospital_death = mean(hospital_death, na.rm = TRUE)) %>%
  ggplot(aes(d1_temp_max, hospital_death)) +
  geom_point()
```



Testing the d1_temp_max on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate,
  data = edx_train)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9226466
```

```
sum(y_hat_rf == 1)
```

```
## [1] 118
```

Testing the d1_temp_max on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heart_rate,
  data = edx_train, method = "C5.0")
```

```
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9224537
```

```
sum(y_hat_ct == 1)
```

```
## [1] 93
```

Testing the d1_temp_max on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate_min)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

```
## [1] 0.9214892
```

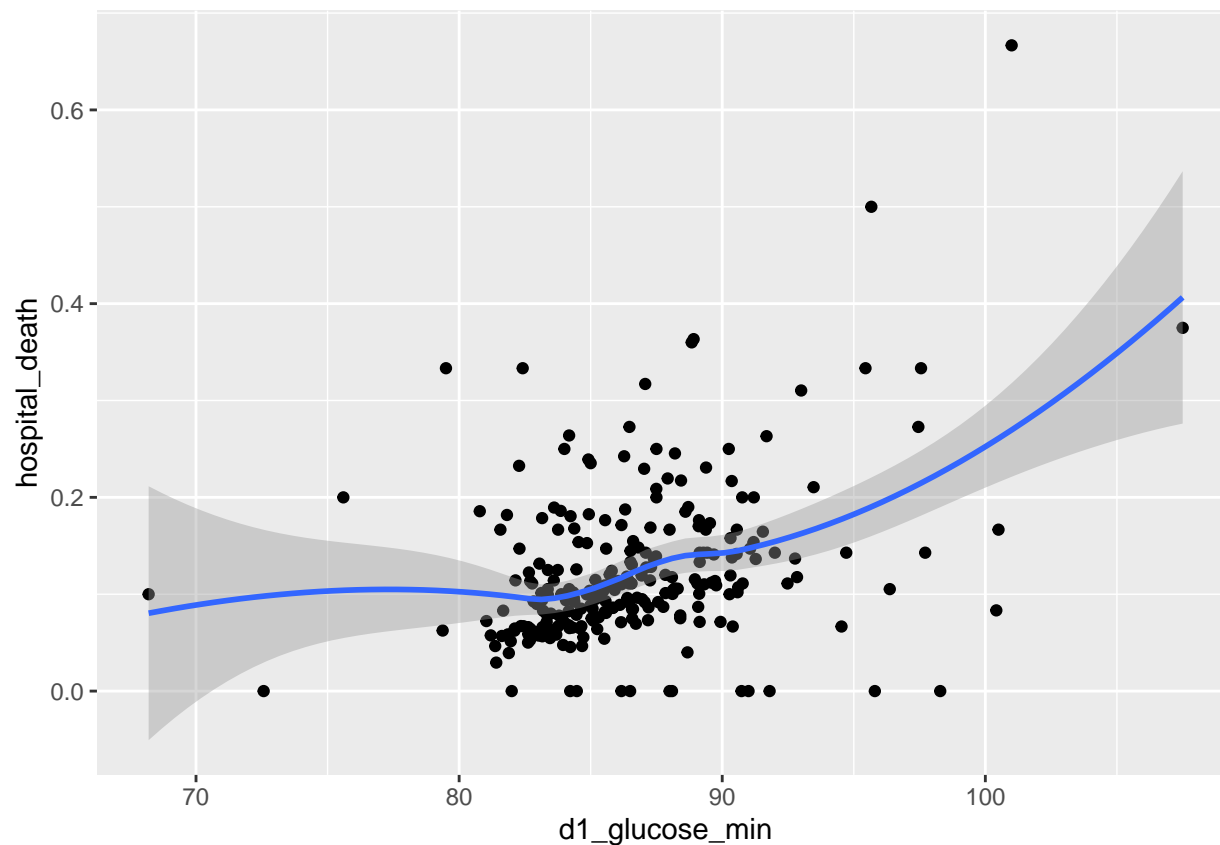
```
sum(y_hat_gam == 1)
```

```
## [1] 94
```

Glucose level is always a factor analyzed by doctors.

```
Survival_Prediction_vi %>%
  group_by(d1_glucose_min) %>%
  summarize(d1_glucose_min = mean(h1_heartrate_min), hospital_death = mean(hospital_death, na.rm = TRUE))
ggplot(aes(d1_glucose_min, hospital_death)) +
  geom_point() +
  geom_smooth()
```

```
## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



There seems to be some relationship between glucose level and hospital death.

Testing the `d1_glucose_min` on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9251543
```

```
sum(y_hat_rf == 1)
```

```
## [1] 111
```

Testing the `d1_glucose_min` on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate)
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9222608
```

```
sum(y_hat_ct == 1)
```

```
## [1] 116
```

Testing the d1__glucose__min on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate + d1__glucose__min)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

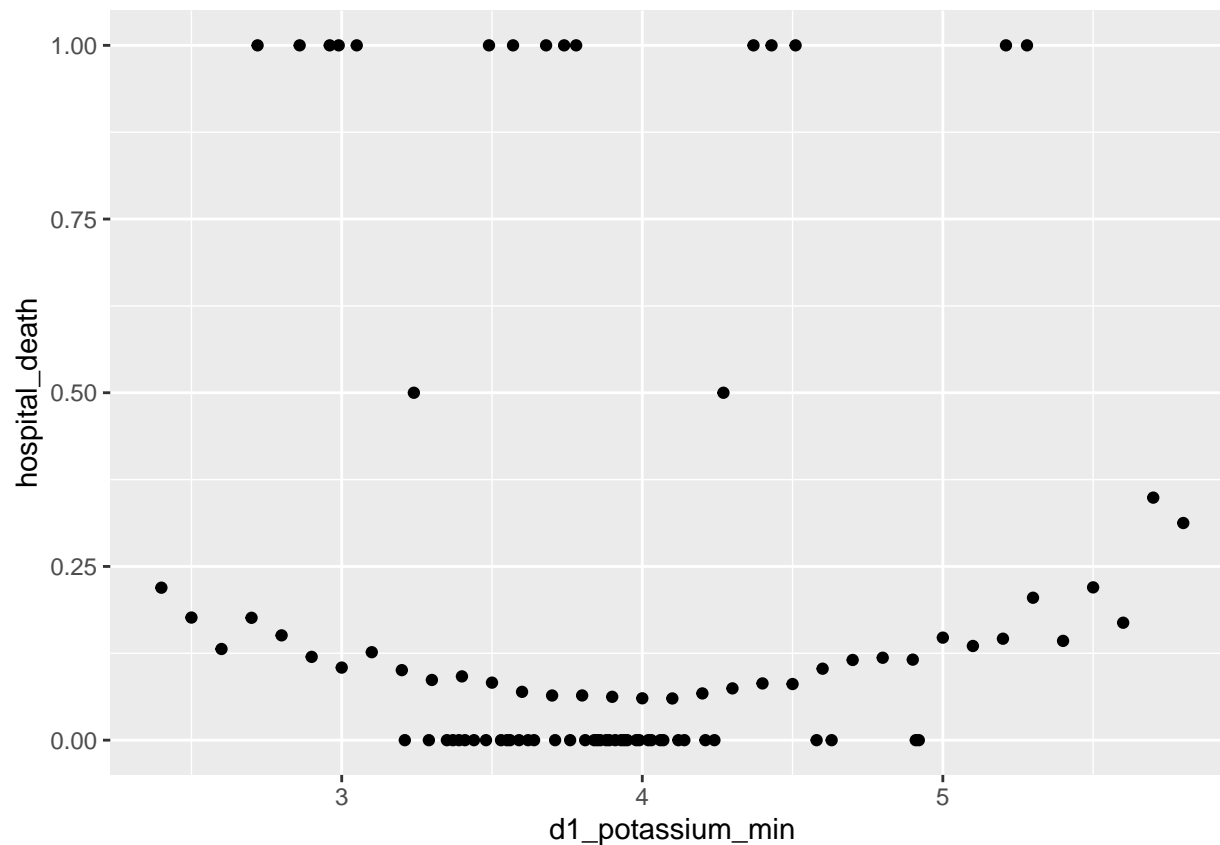
```
## [1] 0.9228395
```

```
sum(y_hat_gam == 1)
```

```
## [1] 97
```

Lack of minerals and poor diet can have an effect on people's health.

```
Survival_Prediction_vi %>%
  group_by(d1_potassium_min ) %>%
  summarize(d1_potassium_min = mean(d1_potassium_min ),hospital_death = mean(hospital_death , na.rm = TRUE))
ggplot(aes(d1_potassium_min ,hospital_death )) +
  geom_point()
```



The points are scattered appearing to be unrelated.

Testing the `d1_potassium_min` on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate,
  data = edx_train,
  ntree = 1000,
  importance = TRUE)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9243827
```

```
sum(y_hat_rf == 1)
```

```
## [1] 107
```

Testing the `d1_potassium_min` on the C5.0Tree algorithm.

```
fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate)
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)
```

```
## [1] 0.9230324
```

```
sum(y_hat_ct == 1)
```

```
## [1] 116
```

Testing the d1_potassium_min on the gamLoess algorithm.

```
fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate + d1_potassium_min)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)
```

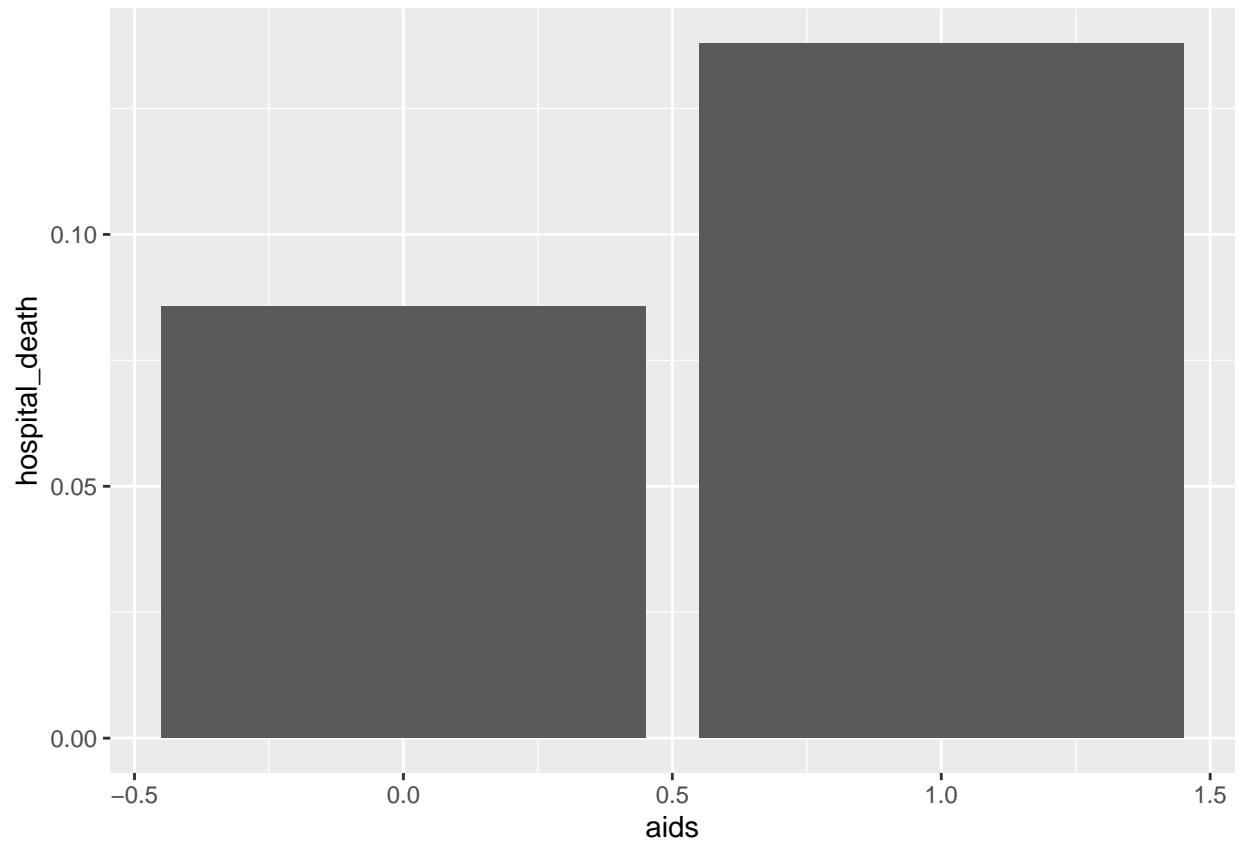
```
## [1] 0.9224537
```

```
sum(y_hat_gam == 1)
```

```
## [1] 101
```

Our intuition says that people who have AIDS are more likely to die. Is this true?

```
Survival_Prediction_vi %>%
  group_by(aids) %>%
  summarize(aids = mean(aids), hospital_death = mean(hospital_death, na.rm = TRUE)) %>%
  ggplot(aes(aids, hospital_death)) +
  geom_col()
```



As we can see people who have AIDS have a little more chance of having a hospital death.

Testing the aids parameter on the random forest algorithm.

```
fit_rf <- randomForest(hospital_death ~ age + ventilated_apache + temp_apache + d1_diasbp_min + d1_heart_rate)

y_hat_rf <- predict(fit_rf, edx_test)

mean(edx_test$hospital_death == y_hat_rf)
```

```
## [1] 0.9232253
```

```
sum(y_hat_rf == 1)
```

```
## [1] 103
```

Testing the aids parameter on the C5.0Tree algorithm.


```

fit_ct <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate)
y_hat_ct <- predict(fit_ct, edx_test)
mean(edx_test$hospital_death == y_hat_ct)

## [1] 0.9230324

sum(y_hat_ct == 1)

## [1] 116

```

Testing the aids parameter on the gamLoess algorithm.

```

fit_gam <- train(hospital_death ~ age + temp_apache + ventilated_apache + d1_diasbp_min + d1_heartrate)
y_hat_gam <- predict(fit_gam, edx_test)
mean(edx_test$hospital_death == y_hat_gam)

## [1] 0.9224537

sum(y_hat_gam == 1)

## [1] 101

```

Results : In this section we are going to use the validation set as our final test and see what algorithm had the best performance and the highest mean.

Testing our final model of the random forest function on the validation set.

```

y_hat_rf <- predict(fit_rf, validation)

mean(validation$hospital_death == y_hat_rf)

## [1] 0.9218886

```

Testing our final model of the C5.0Tree on the validation set.

```
y_hat_ct <- predict(fit_ct, validation)
mean(validation$hospital_death == y_hat_ct)
```

```
## [1] 0.9165076
```

Testing our final model of the gamloess on the validation set.

```
y_hat_gam <- predict(fit_gam, validation)
mean(validation$hospital_death == y_hat_gam)
```

```
## [1] 0.9213678
```

Conclusions :

The random forest got higher mean but its only limitation is that it is very slow to run.

The cs50tree surprisingly doesn't obtain a good perform and a decent mean.

The gamloess algorithm had the second best average of all, but it was very fast to run.

that an algorithm finds a predictor that increases its mean does not mean that in other algorithms it has the same result, for example we have the c50tree and the gamloess.

In a future work we can use more columns of the data table or continue improving the parameters with computers that can support large amounts of information.